

System Programming & OS 실습

10. Concurrency

1

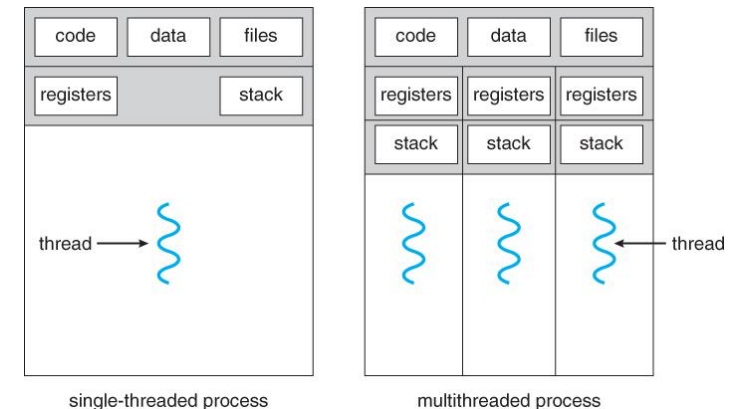
이성현, 최민국

Dankook University

{leesh812, mgchoi}@dankook.ac.kr

- Lock
 - Practice 3: Sloppy Counter
 - Practice 4: Accuracy
- Semaphore
 - Practice 5: Counting semaphore
- Deadlock
 - Practice 6: Dining philosophers problem
 - Practice 7: Circular Wait
 - Practice 8: No Preemption

- Multi-threaded program
 - Thread: flow of control
 - Process: one flow of control + resources (address space, files)
 - Multi-threaded program (or process): multiple flow of controls + resources (address space, files)
 - Multiple threads share address space
 - Multiple Processes do not share their address space
- Concurrency
 - Shared data → race condition → may generate wrong results
 - Concurrency: enforce to access shared data in a synchronized way



(Source: A. Silberschatz, "Operating system Concept")

- Benefit of Thread
 - Fast creation
 - Parallelism
 - Can overlap processing with waiting
 - Data sharing
- Thread management
 - Several stacks in an address space
 - Scheduling entity

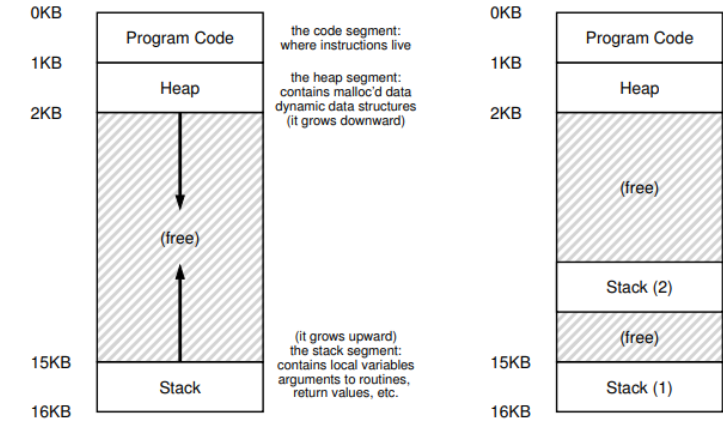
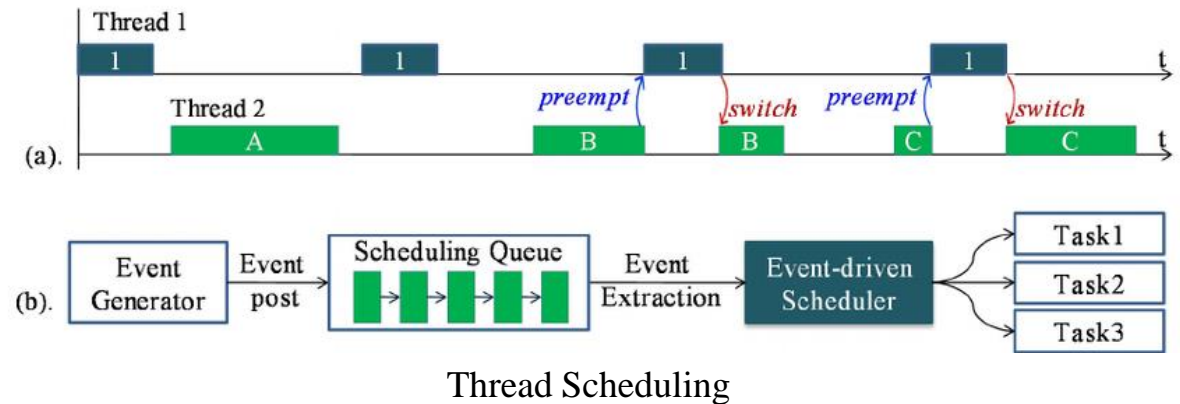


Figure 26.1: Single-Threaded And Multi-Threaded Address Spaces



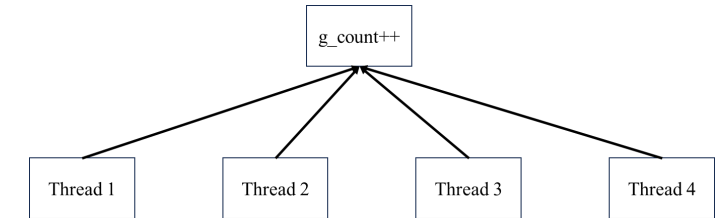
Thread Scheduling

- `#include <pthread.h>`
- `int pthread_create(pthread_t *restrict thread, const pthread_attr_t *restrict attr, void *(*start_routine)(void *), void *restrict arg);`
 - similar to `fork()`, thread exits when the passed function reach the end.
 - `arg1`) thread structure to interact with this thread,
 - `arg2`) attribute of the thread such as priority and stack size, in most case it is `NULL` (use default)
 - `arg3`) function pointer for start routine
 - `arg4`) arguments
- `int pthread_join(pthread_t thread, void **retval);`
 - similar to `wait()`, for synchronization
 - `arg1`) thread structure, which is initialized by the thread creation routine
 - `arg2`) a pointer to the return value (`NULL` means “don’t care”)

• Practice 1

```
> gcc thread.c -lpthread -o thread.out
> ./thread.out 4 1000
> ./thread.out 4 10000
```

(컴파일)
(실행1)
(실행2)



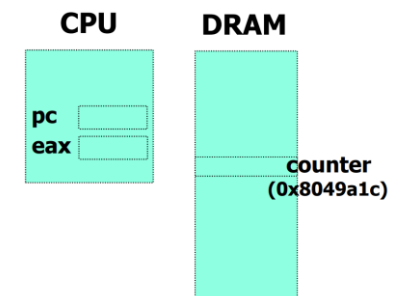
```
embedded@embedded:~/thread_test$ ./thread.out 4 10000
Thread number 0: 10000
Thread number 2: 20000
Thread number 1: 30000
Thread number 3: 40000
Complete
embedded@embedded:~/thread_test$ ./thread.out 4 100000
Thread number 0: 99991
Thread number 2: 218531
Thread number 3: 279583
Thread number 1: 379583
Complete
```

```
17     for (i = 0; i < 1e7; i++) {
18         counter = counter + 1;
19     }
```

```
mov 0x8049a1c, %eax
add $0x1, %eax
mov %eax, 0x8049a1c
```

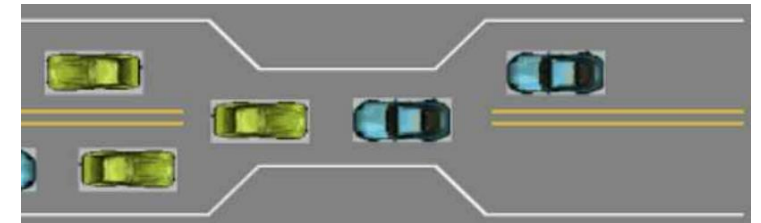
OS	Thread 1	Thread 2	(after instruction)		
			PC	eax	counter
	before critical section		100	0	50
	mov 8049a1c,%eax		105	50	50
	add \$0x1,%eax		108	51	50
interrupt	save T1				
	restore T2		100	0	50
		mov 8049a1c,%eax	105	50	50
		add \$0x1,%eax	108	51	50
		mov %eax,8049a1c	113	51	51
interrupt	save T2				
	restore T1		108	51	51
	mov %eax,8049a1c		113	51	51

Figure 26.7: The Problem: Up Close and Personal



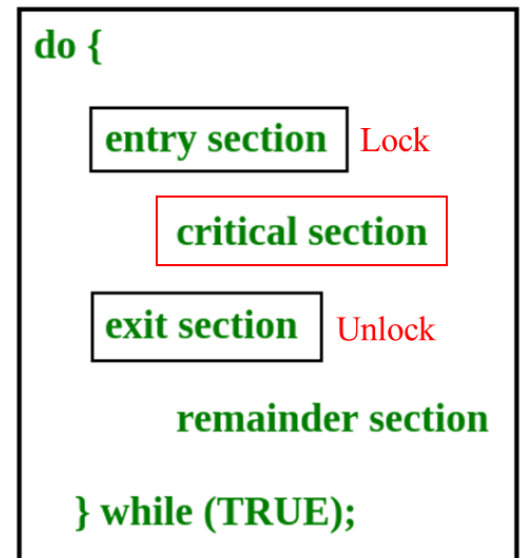
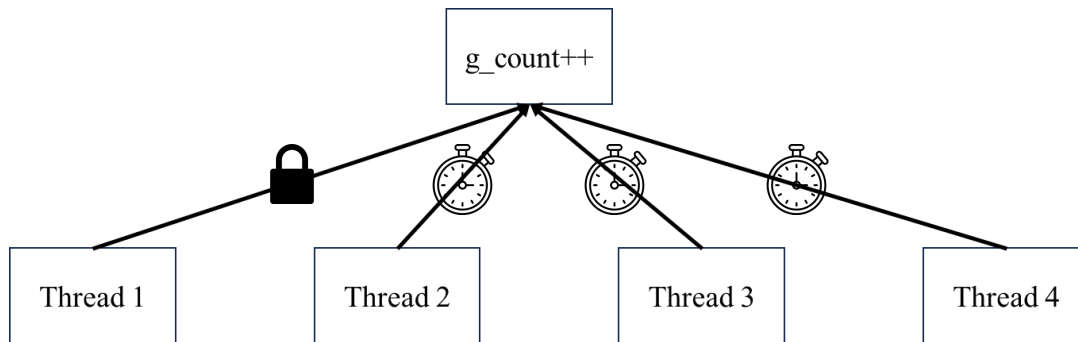
Thread Problem

- Reason
 - Numerous threads access **shared data(critical section)** at the same time
→ **race condition**
 - Uncontrolled scheduling
→ Results are different at each execution depending on scheduling order
- Solution
 - Controlled scheduling: Do all or nothing (indivisible) → **atomicity**
 - The code that can result in the race condition → **critical section**
 - Allow only one thread in the critical section → **mutual exclusion**



Thread Problem

- Mutual exclusion API (mutex_***)
 - #include <pthread.h>
 - pthread_mutex_t lock;
 - int pthread_mutex_init(pthread_mutex_t *restrict mutex,
const pthread_mutexattr_t *restrict attr);
 - int pthread_mutex_lock(pthread_mutex_t *mutex);
 - int pthread_mutex_unlock(pthread_mutex_t *mutex);
 - int pthread_mutex_trylock(pthread_mutex_t *mutex);



- Practice 2

```
> gcc -o thread_lock.out thread_lock.c -lpthread      (컴파일)
> ./thread.out 4 10000                               (실행1)
> ./thread_lock.out 4 10000                         (실행2)
```

```
embedded@embedded:~/thread_test$ ./thread.out 4 100000
```

```
Thread number 0: 99991
```

```
Thread number 2: 218531
```

```
Thread number 3: 279583
```

```
Thread number 1: 379583
```

```
Complete
```

```
embedded@embedded:~/thread_test$ ./thread_lock.out 4 100000
```

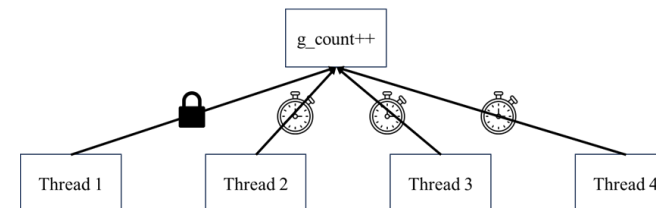
```
Thread number 1: 235328
```

```
Thread number 2: 379740
```

```
Thread number 3: 380224
```

```
Thread number 0: 400000
```

```
Complete
```



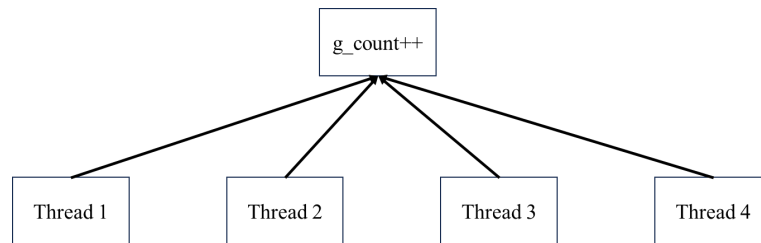
```
do {
    entry section
    critical section
    exit section
    remainder section
} while (TRUE);
```

Practice 3: Sloppy Counter

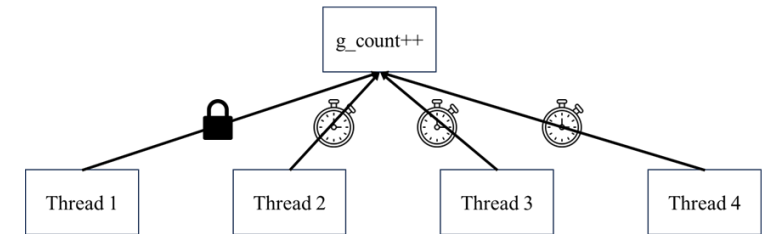
10



Practice 1: Multi-threading without lock



Practice 2: Multi-threading with lock



Which one would be faster?

> time ./thread.out 4 10000000

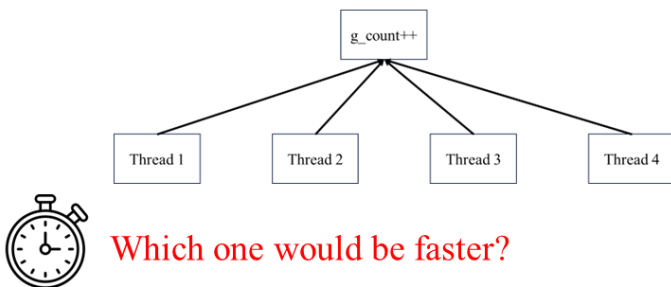
> time ./thread_lock.out 4 10000000

Practice 3: Sloppy Counter

11



Practice 1: Multi-threading without lock



Which one would be faster?

```
> time ./thread.out 4 10000000
```

```
mingu@server:~/TABA_OS_2023/thread_practice$ time ./thread.out 4 10000000
Thread number 2: 10509904
Thread number 1: 12437601
Thread number 0: 12486358
Thread number 3: 13135124
Complete

real    0m0.176s
user    0m0.665s
sys     0m0.000s
```

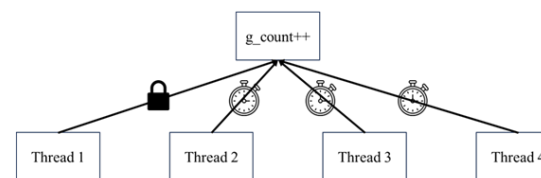


0.176s

14x slower



Practice 2: Multi-threading with lock



```
> time ./thread_lock.out 4 10000000
```

```
mingu@server:~/TABA_OS_2023/thread_practice$ time ./thread_lock.out 4 10000000
Thread number 2: 35600970
Thread number 0: 38511449
Thread number 1: 39049911
Thread number 3: 40000000
Complete

real    0m2.478s
user    0m3.725s
sys     0m4.959s
```



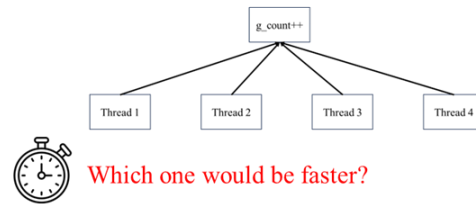
2.478s

Practice 3: Sloppy Counter

12



Practice 1: Multi-threading without lock



> time ./thread.out 4 10000000

```
mingu@server:~/TABA_OS_2023/thread_practice$ time ./thread.out 4 10000000
Thread number 2: 10509904
Thread number 1: 12437601
Thread number 0: 12486358
Thread number 3: 13135124
Complete

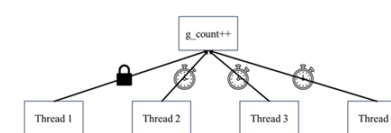
real    0m0.176s
user    0m0.665s
sys     0m0.000s
```



0.176s



Practice 2: Multi-threading with lock



> time ./thread_lock.out 4 10000000

```
mingu@server:~/TABA_OS_2023/thread_practice$ time ./thread_lock.out 4 10000000
Thread number 2: 35600970
Thread number 0: 38511449
Thread number 1: 39049911
Thread number 3: 40000000
Complete

real    0m2.478s
user    0m3.725s
sys     0m4.959s
```



2.478s

14x slower

Execution time becomes too long due to **lock contention** ...

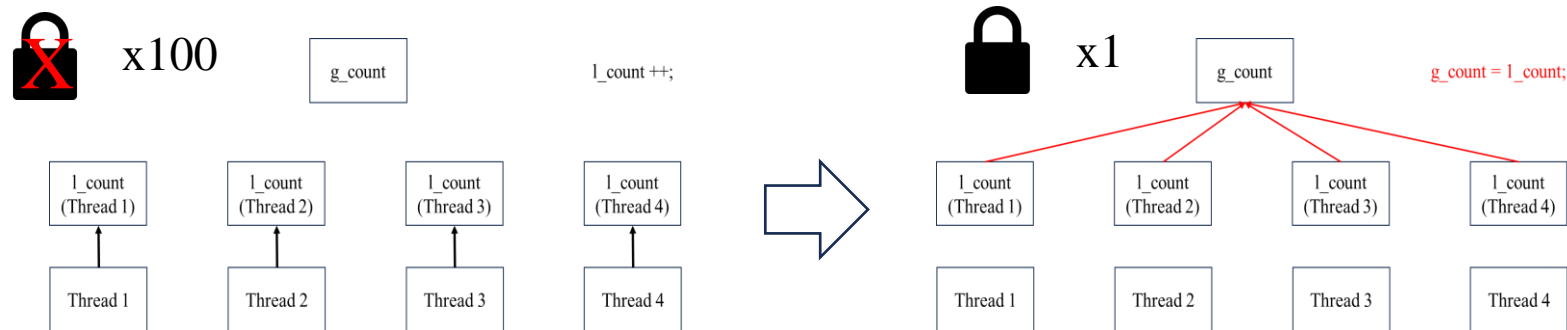
Do we have to lock **g_count** every time?

Or wouldn't it be okay to lock **g_count** just a few times?

Practice 3: Sloppy Counter

13

- Sloppy counter
 - a.k.a Scalable counter or Approximate counter
 - Quite higher performance
 - A single global counter + Several local counters
 - Usually, one per CPU core
 - Update local counter → periodically update global counter
 - Less contention → Scalable



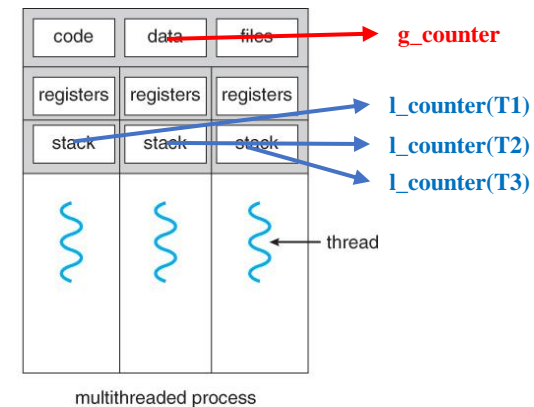
Practice 3: Sloppy Counter

14

- Sloppy counter
 - A single global counter + Several local counters
 - Update local counter → periodically update global counter

Time	L_1	L_2	L_3	L_4	G
0	0	0	0	0	0
1	0	0	1	1	0
2	1	0	2	1	0
3	2	0	3	1	0
4	3	0	3	2	0
5	4	1	3	3	0
6	5 → 0	1	3	4	5 (from L_1)
7	0	2	4	5 → 0	10 (from L_4)

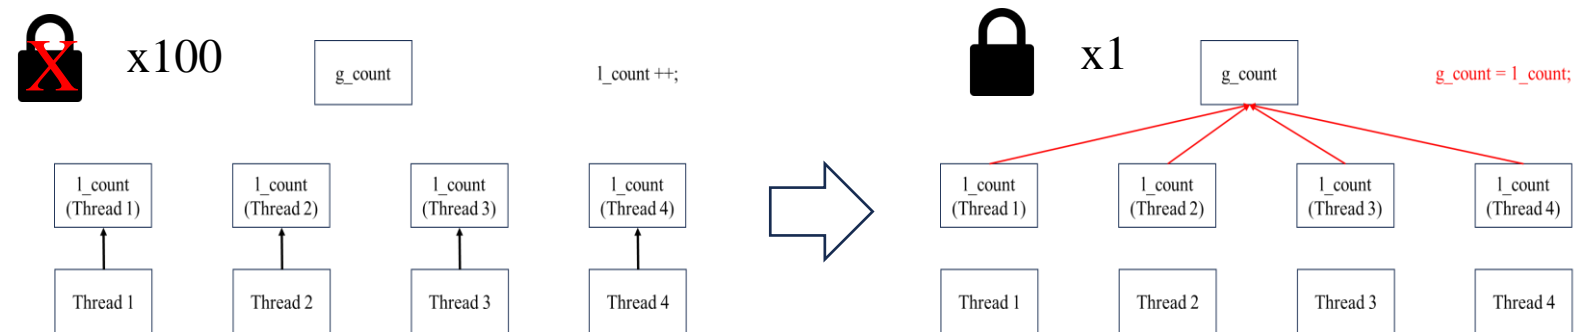
Figure 29.3: Tracing the Approximate Counters



Practice 3: Prepare

15

- Practice 2 command for prepare
 - > cp thread_lock.c thread_sloppy.c (파일 복사)
 - > vim thread_sloppy.c (코드 작성)



```
#include <pthread.h>
#include <stdint.h>

int g_count = 0; // counter (critical section)
int g_nthd = 0; // num of threads
int g_worker_loop_cnt = 0;
int sloppy = 0;
pthread_mutex_t lock;

static void *work(void* tno); // thread routine

int main(int argc, char *argv[]){
    pthread_t *thd_arr; // thread array
    int thd_cnt; // thread count

    if (argc < /*fill the blanks */){
        fprintf(stderr, "%s parameter : nthread, worker_loop_cnt,
sloppy\n", argv[0]);
        exit(-1);
    }

    // get num of threads and worker loop count
    g_nthd = atoi(argv[1]);
    g_worker_loop_cnt = atoi(argv[2]);
    sloppy = /*fill the blanks */;

    // ...
```

```
static void *work(void* cnt){
    int thd_cnt = (int)(intptr_t)cnt;
    int i, j;
    int l_count = 0;

    for(i = 0; i < /*fill the blanks */; i++){
        for(j = 0; j < /*fill the blanks */ ; j++){
            l_count++;
        }
        pthread_mutex_lock(&lock);
        /*fill the blanks */
        /*fill the blanks */

        /*fill the blanks */
    }

    printf("Thread number %d: %d \n", thd_cnt, g_count);
    return NULL;
}
```


- Practice 3 command 2

- gcc -o thread_sloppy.out thread_sloppy.c -lpthread (컴파일)

- > time ./thread_sloppy.out 1 100000 1000

- > time ./thread_sloppy.out 2 100000 1000

- > time ./thread_sloppy.out 3 100000 1000

- > time ./thread_sloppy.out 4 100000 1000

- > time ./thread_lock.out 1 100000

- > time ./thread_lock.out 2 100000

- > time ./thread_lock.out 3 100000

- > time ./thread_lock.out 4 100000

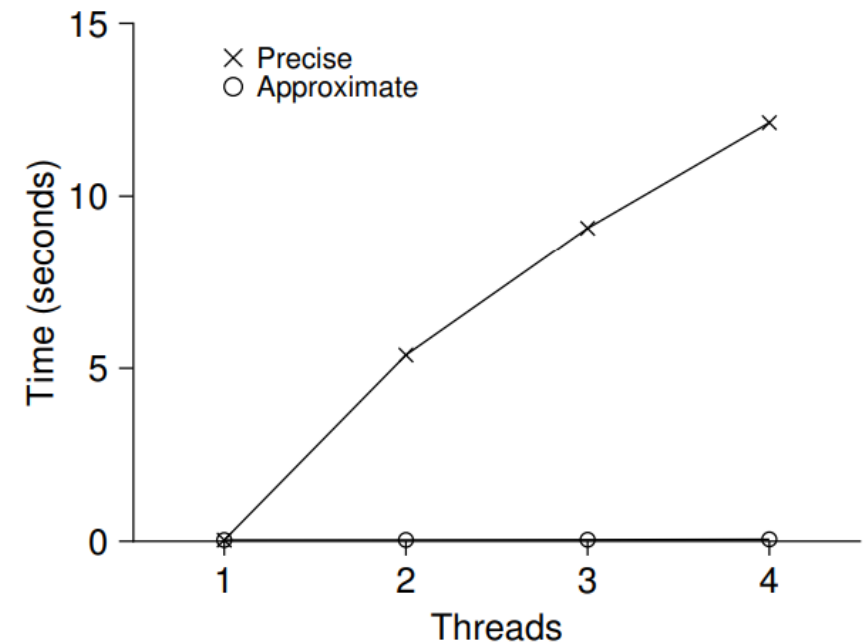


Figure 29.5: Performance of Traditional vs. Approximate Counters

- Practice 3 command 3

- gcc -o thread_sloppy.out thread_sloppy.c -lpthread (컴파일)

- > time ./thread_sloppy.out 4 100000 1
 - > time ./thread_sloppy.out 4 100000 4
 - > time ./thread_sloppy.out 4 100000 16
 - > time ./thread_sloppy.out 4 100000 64
 - > time ./thread_sloppy.out 4 100000 256
 - > time ./thread_sloppy.out 4 100000 1024

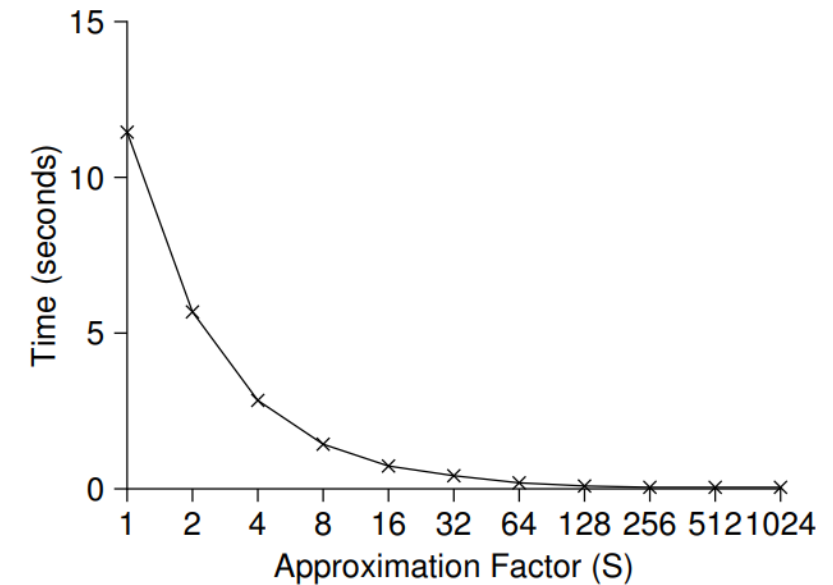


Figure 29.6: Scaling Approximate Counters

Practice 4: Accuracy

19

- Bigger the S , shorter the time
 - Then large S (update frequency) is silver bullet?
- Are there any other considerations?
 - Accuracy of counter during increment

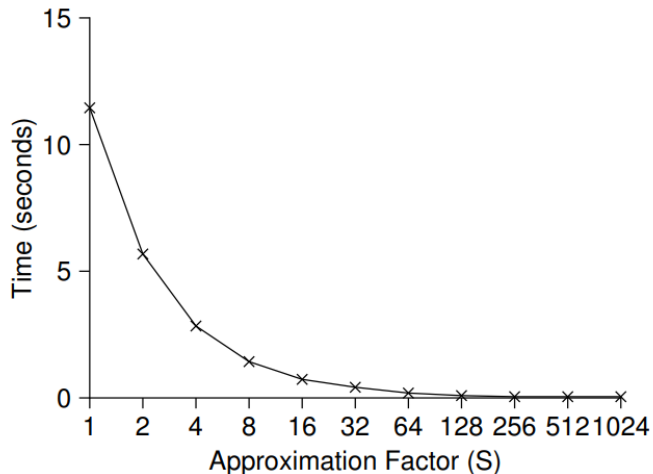


Figure 29.6: Scaling Approximate Counters



Practice 4: Prepare

20

- Practice 4 command for prepare

> cp thread_sloppy.c thread_accuracy.c (파일 복사)

> vim thread_accuracy.c (코드 작성)

Practice 4: Code

21

```
int main(int argc, char *argv[]){
    // ...
    for(thd_cnt=0; thd_cnt < g_nthd; thd_cnt++){
        // create thread
        assert(pthread_create(&thd_arr[thd_cnt], NULL,
            work, (void*) (intptr_t) thd_cnt) == 0);
        pthread_mutex_lock(&lock);
        printf("After create %d: %d \n", thd_cnt, g_count);
        pthread_mutex_unlock(&lock);
    }

    for(thd_cnt=0; thd_cnt < g_nthd; thd_cnt++){
        // join thread
        pthread_mutex_lock(&lock);
        printf("Before join %d: %d \n", thd_cnt, g_count);
        pthread_mutex_unlock(&lock);
        assert(pthread_join(thd_arr[thd_cnt], NULL) == 0);
    }
    printf("Complete\n");
}
```

- Practice 4 command 2

> gcc -o thread_accuracy.out thread_accuracy.c -lpthread (컴파일)

> ./thread_accuracy.out 2 100000 1

> ./thread_accuracy.out 2 100000 10

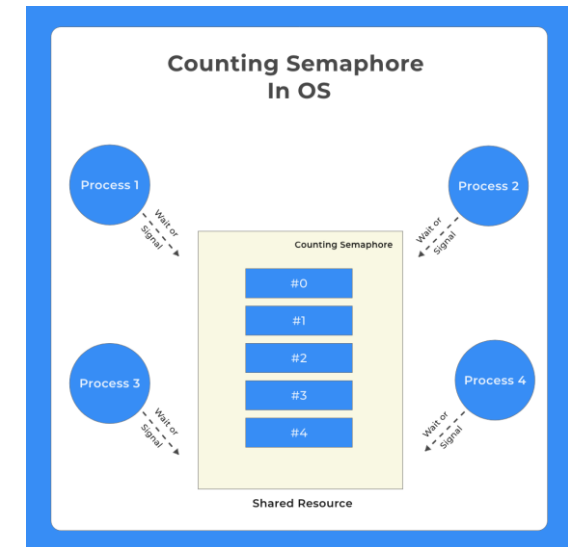
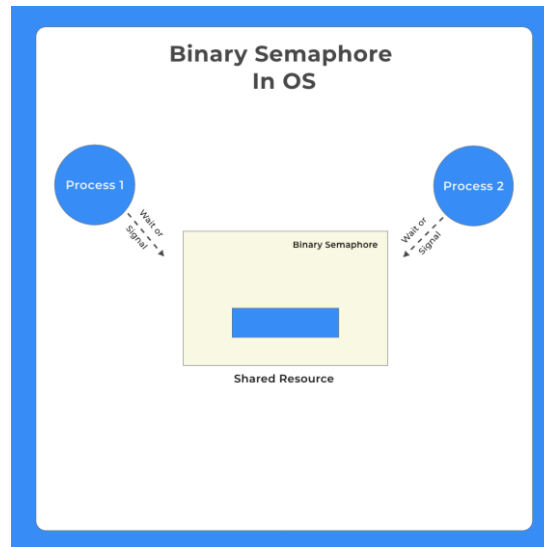
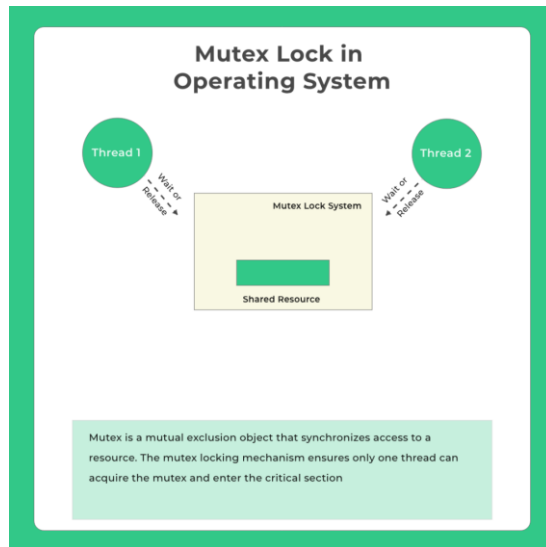
> ./thread_accuracy.out 2 100000 100

> ./thread_accuracy.out 2 100000 1000

```
mingu@server:~/TABA_OS_2023/thread_practice$ ./thread_accuracy.out 2 100000 1
After create 0: 0
After create 1: 906
Before join 0: 1025
Thread number 1: 160644
Thread number 0: 200000
Before join 1: 200000
Complete
mingu@server:~/TABA_OS_2023/thread_practice$ ./thread_accuracy.out 2 100000 10
After create 0: 0
After create 1: 2840
Before join 0: 3810
Thread number 0: 157120
Before join 1: 186960
Thread number 1: 200000
Complete
```

```
mingu@server:~/TABA_OS_2023/thread_practice$ ./thread_accuracy.out 2 100000 100
After create 0: 0
After create 1: 7200
Before join 0: 8200
Thread number 0: 184200
Before join 1: 196100
Thread number 1: 200000
Complete
mingu@server:~/TABA_OS_2023/thread_practice$ ./thread_accuracy.out 2 100000 1000
After create 0: 0
After create 1: 6000
Before join 0: 7000
Thread number 0: 187000
Before join 1: 199000
Thread number 1: 200000
Complete
```

- Semaphore
 - Well-known structure for concurrency control
 - Can be used as both a lock and a condition variable
 - Binary semaphore, Counting semaphore
 - Can be employed by various concurrency problems including
 - 1) producer/consumer, 2) reader/writer and 3) dining philosophers



- Semaphore definition

- An object with an integer value manipulated by three routines

- `sem_init(semaphore, p_shared, initial_value)`

- `sem_wait()`: also called as `P()`, `down()` ...

Decrease the value of the semaphore (S). Then, either return right away (when $S \geq 0$) or cause the caller to suspend execution waiting for a subsequent post (when $S < 0$)

- `sem_post()`: also called as `V()`, `up()`, `sem_signal()` ...

Increment the value of the semaphore and then, if there is a thread waiting to be woken, wakes one of them up

- Others: `sem_trywait()`, `sem_timewait()`, `sem_destroy()`

```
1  #include <semaphore.h>
2  sem_t s;
3  sem_init(&s, 0, 1);
```

Figure 31.1: Initializing A Semaphore

```
1  int sem_wait(sem_t *s) {
2      decrement the value of semaphore s by one
3      wait if value of semaphore s is negative
4  }
5
6  int sem_post(sem_t *s) {
7      increment the value of semaphore s by one
8      if there are one or more threads waiting, wake one
9  }
```

Figure 31.2: Semaphore: Definitions Of Wait And Post

- Using a semaphore as a lock
 - Binary semaphore

```
1  sem_t m;  
2  sem_init(&m, 0, X); // initialize to X; what should X be?  
3  
4  sem_wait(&m);  
5  // critical section here  
6  sem_post(&m);
```

Figure 31.3: A Binary Semaphore (That Is, A Lock)

- Counting semaphore

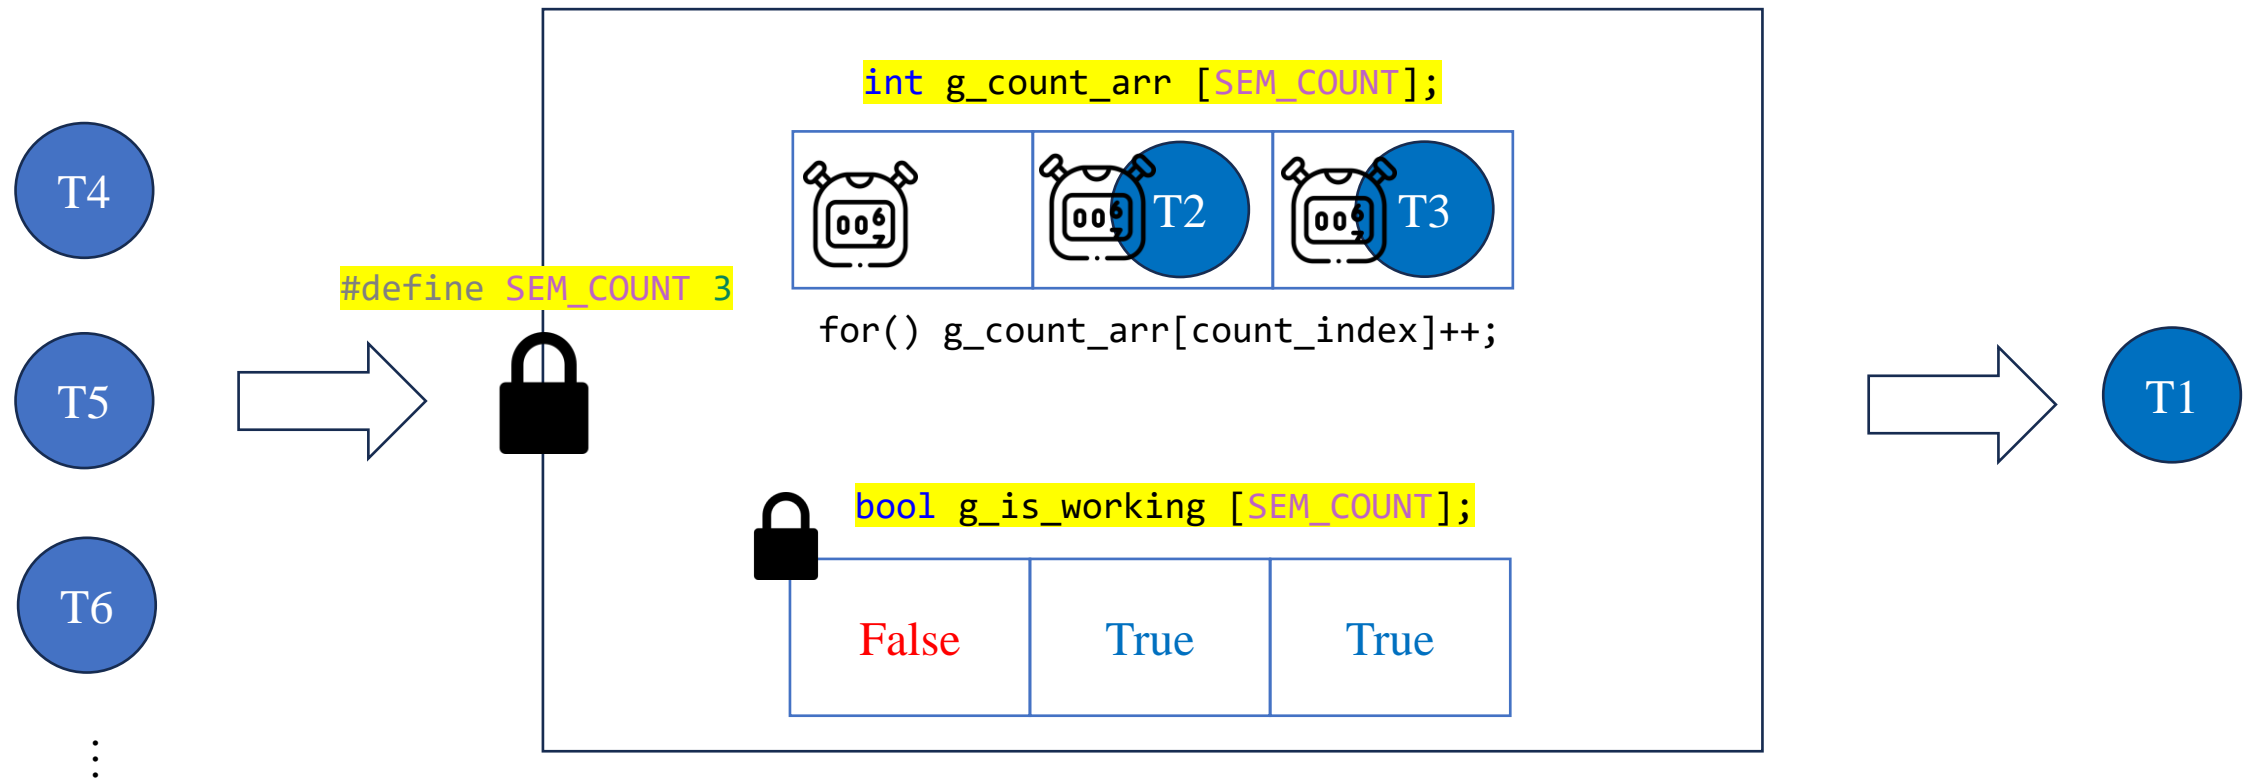
- Using a semaphore as a lock
 - Running example

Val	Thread 0	State	Thread 1	State
1		Run		Ready
1	call sem_wait()	Run		Ready
0	sem_wait() returns	Run		Ready
0	(crit sect begin)	Run		Ready
0	Interrupt; Switch→T1	Ready		Run
0		Ready	call sem_wait()	Run
-1		Ready	decr sem	Run
-1		Ready	(sem<0)→sleep	Sleep
-1		Run	Switch→T0	Sleep
-1	(crit sect end)	Run		Sleep
-1	call sem_post()	Run		Sleep
0	incr sem	Run		Sleep
0	wake(T1)	Run		Ready
0	sem_post() returns	Run		Ready
0	Interrupt; Switch→T1	Ready		Run
0		Ready	sem_wait() returns	Run
0		Ready	(crit sect)	Run
0		Ready	call sem_post()	Run
1		Ready	sem_post() returns	Run

Figure 31.5: Thread Trace: Two Threads Using A Semaphore

Practice 5

27

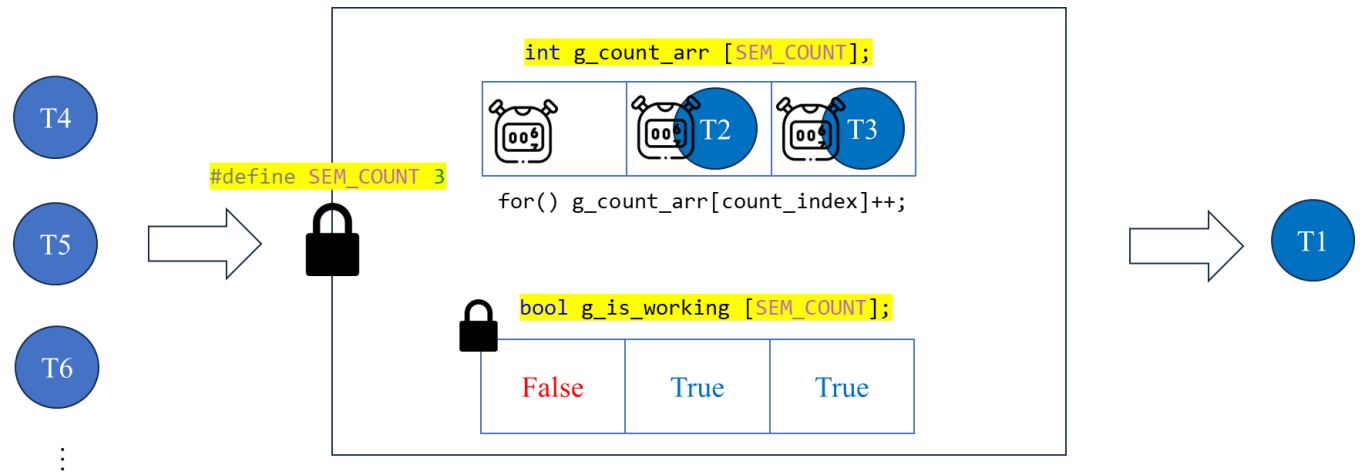


Practice 5: Prepare

28

- Practice 5 command for prepare

> cp thread.c semaphore.c (파일 복사)
> vim semaphore.c (코드 작성)



```
#define SEM_COUNT 3

int g_count = 0; // counter (critical section)
int g_count_arr [SEM_COUNT]; // counter array (critical section)
bool g_is_working [SEM_COUNT]; // check if g_count_arr is working

sem_t b_semaphore; // semaphore for alloc counter index
sem_t c_semaphore; // semaphore for checking # of worker

int main(int argc, char *argv[]){
    // ...
    thd_arr = malloc(sizeof(pthread_t) * g_nthd);

    sem_init(&b_semaphore, /*fill the blanks */); // init sem
    sem_init(&c_semaphore, /*fill the blanks */); // init sem

    // ...

    for(thd_cnt=0; thd_cnt<g_nthd; thd_cnt++) {
        assert(pthread_join(thd_arr[thd_cnt], NULL) == 0);
    }

    // check result of counter array
    printf("Counter array: ");
    for (int i = 0; i < SEM_COUNT; i++){
        printf("%d\t", g_count_arr[i]);
    }
    printf("\n");
}
```

```
static void *work (void* cnt){
    int thd_cnt = (int)(intptr_t)cnt;
    int count_index = -1; // index for count arr

    sem_wait(/*fill the blanks */);

    sem_wait(/*fill the blanks */);
    for (int i=0; i < SEM_COUNT; i++){
        if(g_is_working[i] == false){ // check if counter is not working
            g_is_working[i] = /*fill the blanks */; // this counter will be used
            count_index = /*fill the blanks */; // remember counter index
            break;
        }
    }
    sem_post(/*fill the blanks */);

    if(count_index == -1) {
        fprintf(stderr, "Thread number %d: count_index < 0", thd_cnt);
        exit(-1);
    }

    for(int i = 0; i < g_worker_loop_cnt; i++)
        g_count_arr[count_index]++;

    sem_wait(/*fill the blanks */);
    g_is_working[count_index] = /*fill the blanks */; // free counter
    sem_post(/*fill the blanks */);

    sem_post(/*fill the blanks */);

    return NULL;
}
```

- Command

> gcc -o semaphore.out semaphore.c -lpthread (컴파일)

> ./semaphore.out 10 10

> ./semaphore.out 10 100000

```
mingu@server:~/TABA_OS_2023/thread_practice$ gcc -o semaphore.out semaphore.c -lpthread
mingu@server:~/TABA_OS_2023/thread_practice$ ./semaphore.out 10 10
Count array: 100      0      0
mingu@server:~/TABA_OS_2023/thread_practice$ ./semaphore.out 10 100000
Count array: 400000    300000  300000
mingu@server:~/TABA_OS_2023/thread_practice$
```

- Deadlock

- A situation where two or more threads wait for events that never occur

```
Thread 1:  
pthread_mutex_lock (L1);  
pthread_mutex_lock (L2);
```

```
Thread 2:  
pthread_mutex_lock (L2);  
pthread_mutex_lock (L1);
```

- E.g.) When a thread (say Thread 1) is holding a lock (L1) and waiting for another one (L2); unfortunately, the thread (Thread 2) that holds lock L2 is waiting for L1 to be released.

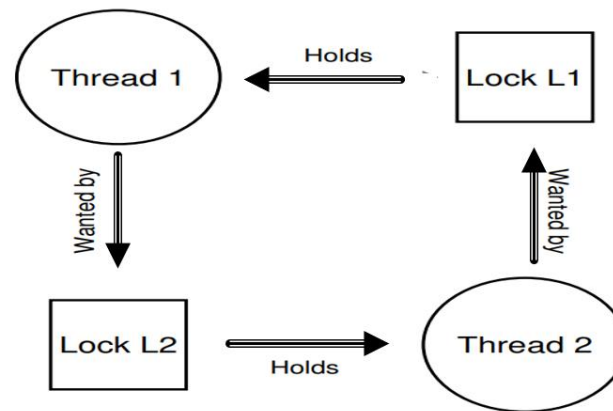


Figure 32.2: The Deadlock Dependency Graph

- Deadlock: 4 Conditions
 - Mutual exclusion
 - Hold-and-Wait
 - No preemption for resource
 - Circular wait

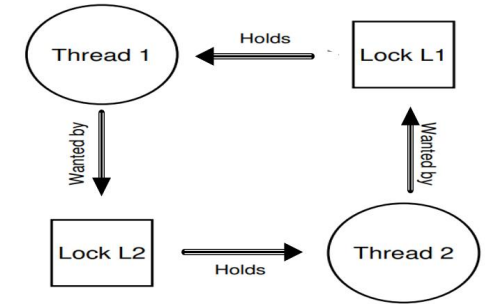
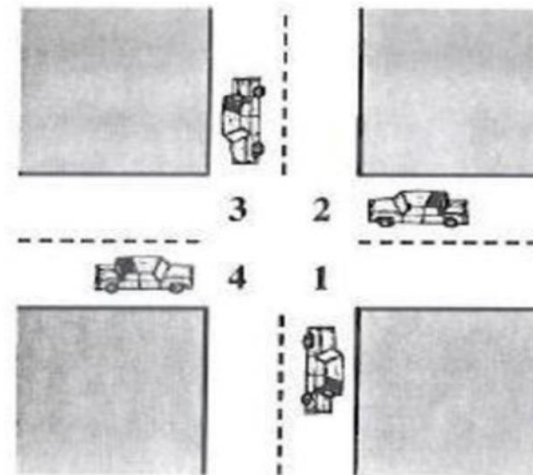
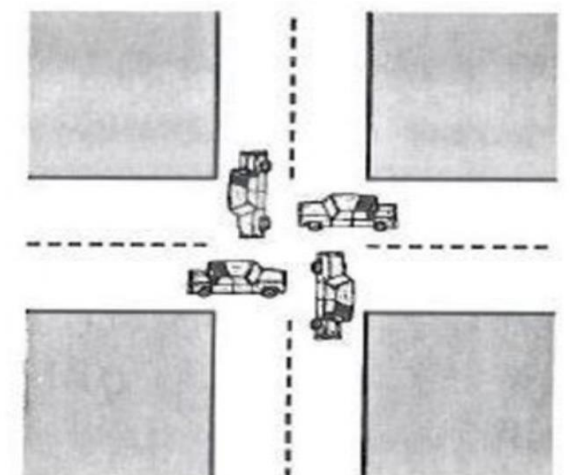


Figure 32.2: The Deadlock Dependency Graph



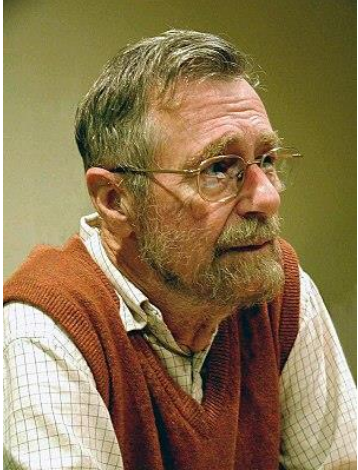
(a) Deadlock Possible



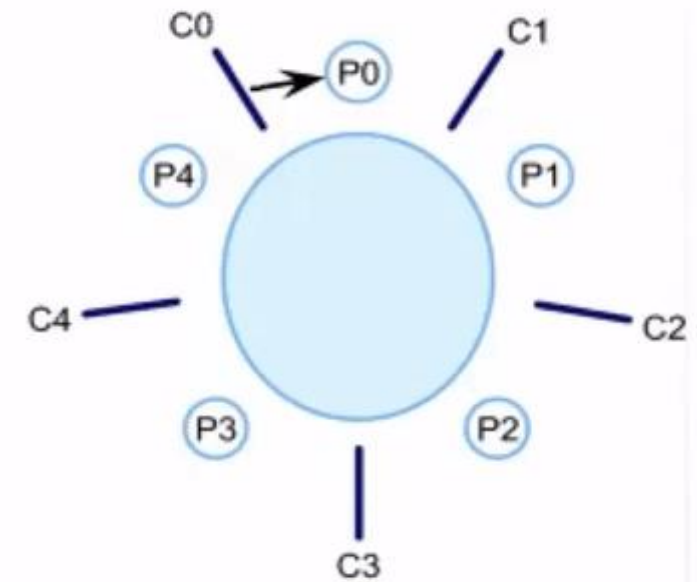
(b) Deadlock

Dining philosophers problem

33



Edsger Wybe Dijkstra



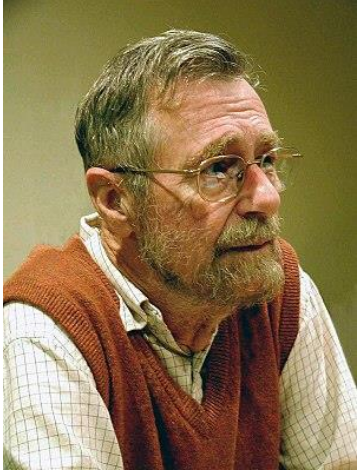
P = Philosopher
C = Chopstick

Problem Definition

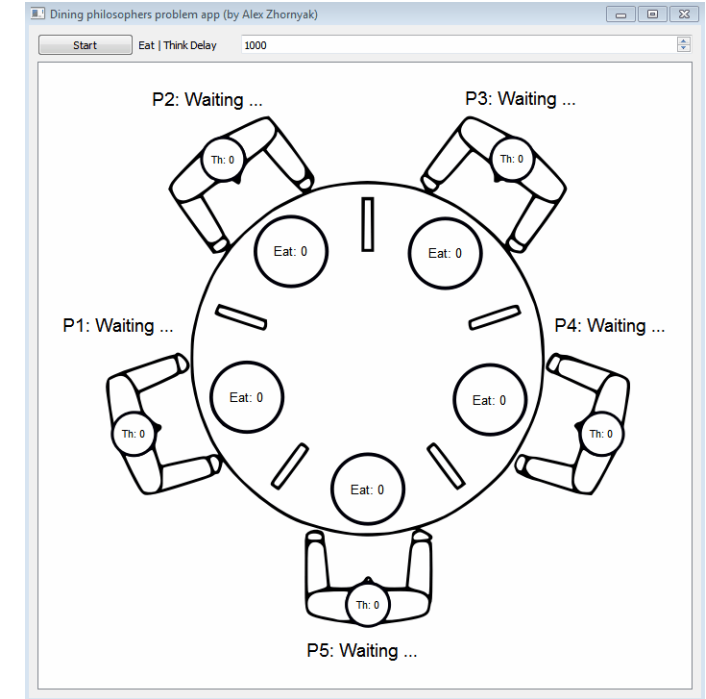
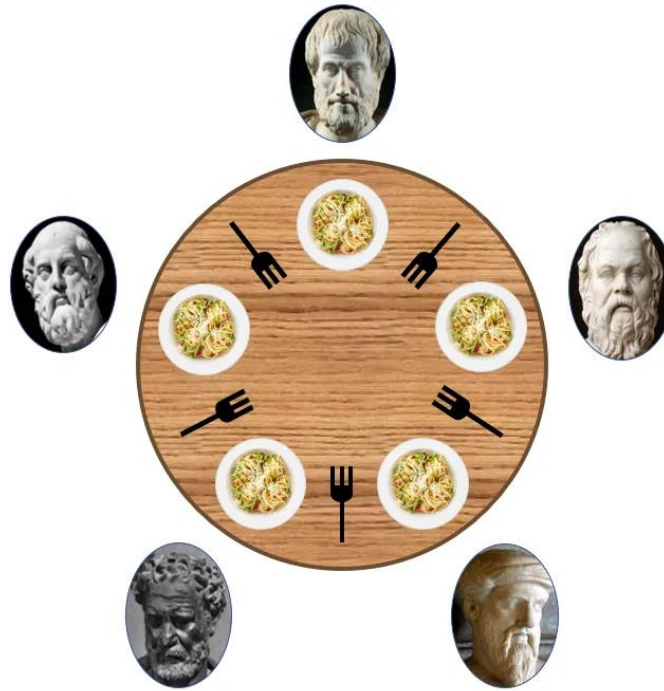
- There are five “philosophers” sitting around a table.
- Between each pair of philosophers is a single fork (thus, five total)
- The philosophers each have times for thinking or for eating
- In order to eat, a philosopher needs two forks, both the one on their left and the one on their right → shared resource → concurrency

Dining philosophers problem

34



Edsger Wybe Dijkstra



P6: Dining philosophers problem

35

- Prepare
 - > cd ~
 - > mkdir philosophers
 - > cd philosophers
 - > vim philosophers.c

P6: Dining philosophers problem

36

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

#define NUM_PHILS 5
pthread_mutex_t forks[NUM_PHILS];

void init();
int leftOf(int i);
int rightOf(int i);
void* philosopher(void* param);
void think(int id);
void eat(int id);
void get_forks(int id);
void put_forks(int id);

int main(){
    pthread_t *thd_arr; // thread array
    thd_arr = malloc(sizeof(pthread_t) * NUM_PHILS);

    for(int i = 0; i < NUM_PHILS; i++){
        pthread_mutex_init(&forks[i], NULL);
    }

    for(int i = 0; i < NUM_PHILS; i++){
        pthread_create(&thd_arr[i], NULL,
                      philosopher, (void*) &i);
        usleep((1 + rand() % 50) * 10000);
    }

    for(int i = 0; i < NUM_PHILS; i++){
        pthread_join(thd_arr[i], NULL);
    }
    return 0;
}

int leftOf(int i){
    return (i) % NUM_PHILS;
}

int rightOf(int i){
    return (i + 1) % NUM_PHILS;
}

void* philosopher(void* param){
    int id = *((int *) param);
    while(1){
        think(id);
        get_forks(id);
        eat(id);
        put_forks(id);
    }
}

void think(int id){
    printf("%d: Now, I'm thinking...\n", id);
}

void eat(int id){
    printf("%d: Now, I'm eating...\n", id);
}

void get_forks(int id){
    pthread_mutex_lock(&forks[leftOf(id)]);
    pthread_mutex_lock(&forks[rightOf(id)]);
}

void put_forks(int id){
    pthread_mutex_unlock(&forks[leftOf(id)]);
    pthread_mutex_unlock(&forks[rightOf(id)]);
}
```

P6: Dining philosophers problem

37

- Run

- > gcc -o philosophers.out philosophers.c -lpthread

- > ./philosophers.out

```
mingu@server:~/TABA_OS_2023/philosophers$ gcc -o philosophers philosophers.c -lpthread
mingu@server:~/TABA_OS_2023/philosophers$ ./philosophers
1: Now, I'm thinking...
3: Now, I'm thinking...
0: Now, I'm thinking...
2: Now, I'm eating...
2: Now, I'm thinking...
2: Now, I'm eating...
2: Now, I'm thinking...
1: Now, I'm eating...
1: Now, I'm thinking...
0: Now, I'm eating...
3: Now, I'm eating...
3: Now, I'm thinking...
4: Now, I'm thinking...
4: Now, I'm eating...
0: Now, I'm thinking...
2: Now, I'm eating...
4: Now, I'm thinking...
1: Now, I'm eating...
2: Now, I'm thinking...
0: Now, I'm eating...
1: Now, I'm thinking...
0: Now, I'm thinking...
```

Deadlock

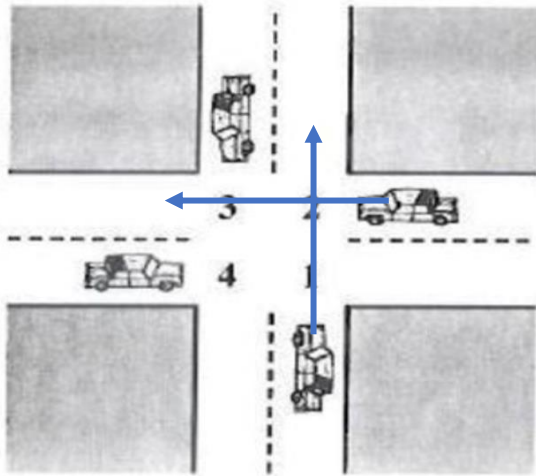
- How to handle Deadlock: three strategies
 - 1. Deadlock Prevention
 - 2. Deadlock Avoidance via Scheduling
 - 3. Deadlock Detection and Recovery

- Deadlock prevention
 - This strategy seeks to prevent one of the 4 Deadlock conditions
 - 1. Hold-and-wait
 - Acquire all locks at once, atomically

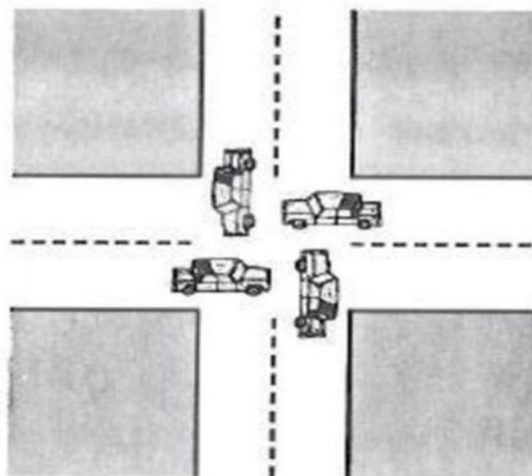
```
1  pthread_mutex_lock(prevention); // begin lock acquisition
2  pthread_mutex_lock(L1);
3  pthread_mutex_lock(L2);
4  ...
5  pthread_mutex_unlock(prevention); // end
```

Acquire all locks atomically (Super-Lock)

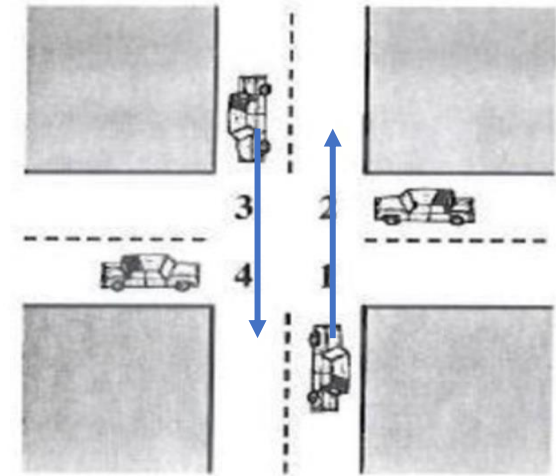
- Deadlock prevention
 - This strategy seeks to prevent one of the 4 Deadlock conditions
 - 2. Circular Wait
 - A total ordering on lock acquisition



(a) Deadlock Possible



(b) Deadlock



(a) Deadlock Possible

- Prepare
 - > cp philosophers.c cp philosophers_circular.c
 - > vim philosophers_circular.c

P7: Circular Wait

42

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

#define NUM_PHILS 5
pthread_mutex_t forks[NUM_PHILS];

void init();
int leftOf(int i);
int rightOf(int i);
void* philosopher(void* param);
void think(int id);
void eat(int id);
void get_forks(int id);
void put_forks(int id);

int main(){
    pthread_t *thd_arr; // thread array
    thd_arr = malloc(sizeof(pthread_t) * NUM_PHILS);

    for(int i = 0; i < NUM_PHILS; i++){
        pthread_mutex_init(&forks[i], NULL);
    }

    for(int i = 0; i < NUM_PHILS; i++){
        pthread_create(&thd_arr[i], NULL,
                      philosopher, (void*) &i);
        usleep((1 + rand() % 50) * 10000);
    }

    for(int i = 0; i < NUM_PHILS; i++){
        pthread_join(thd_arr[i], NULL);
    }
    return 0;
}
```

```
int leftOf(int i){
    return (i) % NUM_PHILS;
}

int rightOf(int i){
    return (i + 1) % NUM_PHILS;
}

void* philosopher(void* param){
    int id = *((int *) param);
    while(1){
        think(id);
        get_forks(id);
        eat(id);
        put_forks(id);
    }
}

void think(int id){
    printf("%d: Now, I'm thinking...\n", id);
}

void eat(int id){
    printf("%d: Now, I'm eating...\n", id);
}
```

```
void get_forks(int id){
    if (forks[leftOf(id)] == 0){
        forks[leftOf(id)] = 1;
    } else{
        pthread_mutex_lock(&forks[leftOf(id)]);
        pthread_mutex_lock(&forks[rightOf(id)]);
    }
}

void put_forks(int id){
    pthread_mutex_unlock(&forks[leftOf(id)]);
    pthread_mutex_unlock(&forks[rightOf(id)]);
}
```

- Run

- > gcc -o philosophers_circular.out philosophers_circular.c -lpthread

- > ./ philosophers_circular.out

- Deadlock prevention

- This strategy seeks to prevent one of the 4 Deadlock conditions

- 3. No Preemption

- Release lock if it can not hold another lock

- Concern: 1) may cause Livelock, 2) sometimes require undo

- • Two threads could both be repeatedly attempting this sequence and repeatedly failing to acquire both locks → add random delay

```
1  top:
2      pthread_mutex_lock(L1);
3      if (pthread_mutex_trylock(L2) != 0) {
4          pthread_mutex_unlock(L1);
5          goto top;
6      }
```

Release lock if it can not hold another lock

P8: No Preemption

45

- Prepare
 - > cp philosophers.c cp philosophers_no_preemphion.c
 - > vim philosophers_no_preemphion.c

P8: No Preemption

46

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

#define NUM_PHILS 5
pthread_mutex_t forks[NUM_PHILS];

void init();
int leftOf(int i);
int rightOf(int i);
void* philosopher(void* param);
void think(int id);
void eat(int id);
void get_forks(int id);
void put_forks(int id);

int main(){
    pthread_t *thd_arr; // thread array
    thd_arr = malloc(sizeof(pthread_t) * NUM_PHILS);

    for(int i = 0; i < NUM_PHILS; i++){
        pthread_mutex_init(&forks[i], NULL);
    }

    for(int i = 0; i < NUM_PHILS; i++){
        pthread_create(&thd_arr[i], NULL,
            philosopher, (void*) &i);
        usleep((1 + rand() % 50) * 10000);
    }

    for(int i = 0; i < NUM_PHILS; i++){
        pthread_join(thd_arr[i], NULL);
    }
    return 0;
}
```

```
int leftOf(int i){
    return (i) % NUM_PHILS;
}

int rightOf(int i){
    return (i + 1) % NUM_PHILS;
}

void* philosopher(void* param){
    int id = *((int *) param);
    while(1){
        think(id);
        get_forks(id);
        eat(id);
        put_forks(id);
    }
}

void think(int id){
    printf("%d: Now, I'm thinking...\n", id);
}

void eat(int id){
    printf("%d: Now, I'm eating...\n", id);
}
```

```
void get_forks(int id){
    while (1){
        /*fill blanks*/
        if (/*fill blanks*/ != 0){
            /*fill blanks*/
        } else{
            /*fill blanks*/
        }
    }
}

void put_forks(int id){
    pthread_mutex_unlock(&forks[id]);
    pthread_mutex_unlock(&forks[id]);
}
```

- Run

- > gcc -o philosophers_no_preemption.out philosophers_no_preemption.c -lpthread

- > ./ philosophers_no_preemption.out

- Deadlock prevention

- This strategy seeks to prevent one of the 4 Deadlock conditions

- 4. Mutual Exclusion:

- “lock free” approach: no lock but support mutual exclusion

- Using powerful hardware instructions, we can build data structures in a manner that does not require explicit locking

- Atomic integer operation with compare-and-swap (chapter 28.9 in LN 4)

```
void increment(counter_t *c) {  
    Pthread_mutex_lock(&c->lock);  
    c->value++;  
    Pthread_mutex_unlock(&c->lock);  
}
```

Using Lock

```
1 void AtomicIncrement(int *value, int amount) {  
2     do {  
3         int old = *value;  
4     } while (CompareAndSwap(value, old, old + amount) == 0);  
5 }
```

Lock free

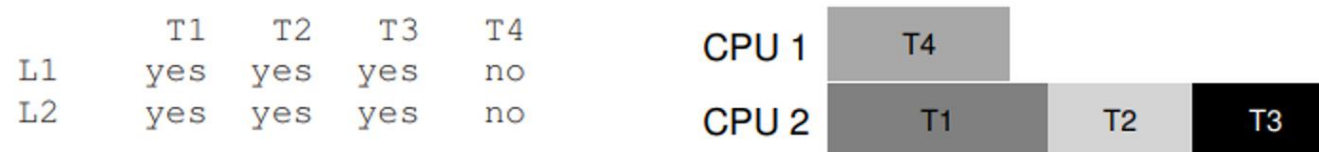
- Deadlock Avoidance via Scheduling

- Instead of prevention, try to avoid by scheduling threads in a way as to guarantee no deadlock can occur.

- E.g.) two CPUs, four threads, T1 wants to use L1 and L2, T2 also wants both, T3 wants L1 only, T4 wants nothing



- E.g. 2) more contention (negative for load balancing)



- No deadlock, but under-utilization → A conservative approach